

Lesson 4: Instance Based Learning

Instance based learning

Instance based learning distinguishes machine learning models which perform “**lazy learning**” at **scoring-time directly from the training data**. This is different from **other** machine learning models which lossily compress the training data into a **simpler hypothesis function**.

Some benefits

- Remembers all training data points.
- Fast for known points (lookup – depending on properties of the data structure used to store).
- “Simple”

Possible problems

- Overfits and generalizes badly
- Must handle conflicting training data

Since instance based learning often requires **traversing the training data** itself, it may be slow to score unknown inputs, but is easily adaptable to new training data points.

k-Nearest Neighbors

The number of neighbors k is a free parameter which must be learned.

Given: * training data $D = \{x_i, y_i\}$, * distance metric $d(q, x)$ * the required number of neighbors k * the query point q :

$$NN = \{i : d(q, x_i) \text{ in } k \text{ smallest}\}$$

For **classification**: vote of y_i or weighted vote or some other strategy.

For **regression**: take mean or weighted mean (weights may be $1/\text{distance}$)

If there are distance conflicts (exactly same distances), “**take all of them**” the way college rankings do (this fudges the k).

Comparing learning and query times

Given a list of n data points (x_i, y_i) sorted by x_i what is the running time and space consumption of “learning” (assume no effort for ETL – which make no sense) and “query” (do not consider the learned data structure’s space consumption) of 1-NN, k -NN, and linear regression:

model	Running Time	Space
1-NN Learning	1	n
1-NN Querying	$\log n$ (binary search)	1
k-NN Learning	1	n
k-NN Querying	$\log n + k$	1
Linear Regression Learning	n	1
Linear Regression Querying	1	1

- Simple linear regression learning runs in $O(n)$ (due to bound on the dimension of x_i , even normal equations solution is “linear” – though this is a bit misleading) and takes $O(1)$ space (For parameter).

Querying in instance based algorithms may be slower (depending on the data structure), though assuming no ETL costs, learning in instance based algorithms is faster. Moreover, **incremental learning** of new data points are possible and probably very fast.

IBL is more lazy about learning (pushes it till querying). So k-NN is a **lazy learner** (or “just in time” learner) whereas linear regression is an **eager learner**.

K-NN Bias

Preference bias: when searching the hypothesis space, the search may “prefer” a particular subset of the hypothesis space over another. Perhaps the hypothesis space is not completely searched and rather only a subset is searched. E.g. simpler trees, simpler functions, Occam’s razor. Compare to restriction bias which is the total representational power of the hypothesis space itself.

What about k-NN?

- **Locality** – near points are similar and how is nearness defined. Moreover, define distance on input not label/output. This comes from domain knowledge.
- **Smoothness** – averaging produces smoothness as opposed to discontinuities.
- All **features seem to matter “equally”** – this comes from the distance function too.

Curse of dimensionality

As the number of features or dimensions grows, the **amount of data we need** to generalize accurately **grows exponentially**. Comes from Richard Bellman, the dynamic programming guy.

Covering the “same space” in higher dimensions requires exponentially more points. (“Same space” is not well-defined.)

Applies in all of ML and not just k-NN.

Distance function weighted on the dimensionality may be a solution.

Some other points

Distance metrics $d(x, q)$ – e.g. Euclidean, Manhattan, weighted-components, mismatches (for classification), etc. But the choice really matters.

Implications of k .

- When $n = k$, we get a constant function (assuming vote or regular average).
- Let $k = n$ with weighted average, the computation may not be so easy, but “smooth”-ish hypothesis. Points closer have greater say.
- Perhaps instead of weighted average, do “local” regression on the k -closest-points. Known as “**locally weighted regression**”. We can do decision trees, neural nets, etc. there. We can achieve more sophisticated curves and more complex hypotheses (removes some restriction bias, but may overfit).

What have we learned?

- Instance based learning
- Eager and lazy learning – lazy puts off work until needed, eager generally learns from training data ahead of time.
- k-NN
- nearest-neighbor; similarity (distance)
- Domain knowledge matters (distance)
- Classification versus regression
- “Averaging”/combining results.
- Composing learning algorithms together (locally weighted regression & regression).
- Curse of dimensionality. Required data is $O(2^d)$.
- No free lunch theorem(?) – averaging over all possible data makes any learning algorithm no better than random. Domain knowledge helps.